

Recurrent Reasoning Networks for the Maximum Independent Set Problem

Malik Mardan

Matriculation number 3834310 Study programme: Master Computer Science

July 1, 2026

This project adapts the Tiny Recursive Model (TRM) of Jolicoeur-Martineau et al., *Less is More: Recursive Reasoning with Tiny Networks* (arXiv:2510.04871), from grid puzzles to graph problems. The idea of TRM is that a very small network, applied many times in a recurrent loop, can solve hard reasoning tasks that usually call for much larger models. Code: <https://github.com/73gon/trm-mis>.

1 Approach

Problem. Given an undirected graph $G = (V, E)$, the Maximum Independent Set (MIS) problem asks for the largest set $S \subseteq V$ in which no two nodes share an edge. MIS is NP-hard. We frame it as per-node classification: the network outputs a probability $p_v \in [0, 1]$ that node v belongs to an optimal set, and a deterministic decoder turns these probabilities into a valid independent set.

Architecture. The model (**GraphTransformerTRM**, about 1.55 M parameters) combines a graph-transformer backbone with the recurrent reasoning loop of TRM:

- **Inputs.** Each node carries 12 local structural features (degree and normalised degree, clustering coefficient, neighbour-degree statistics, and k -core number) plus a 16-dimensional positional encoding. These are embedded and concatenated into a 256-dimensional vector.
- **GPS layers** (2 layers, hidden width 256, 4 attention heads). Every layer runs a local GIN message-passing branch and a global self-attention branch in parallel, so a node sees both its neighbourhood and the whole graph.
- **Recurrent reasoning.** A latent state z and an output state y are refined by $H=2$ outer cycles of $L=6$ inner updates, giving 12 reasoning steps in total. Only the final step backpropagates; the earlier ones run without gradients. This lets the small network improve its answer by repeating computation instead of adding parameters, which is the central claim of TRM.

Training. The model is trained with full supervision. The loss adds a class-weighted binary cross-entropy on the node labels to a soft feasibility penalty that punishes giving high probability to both endpoints of an edge:

$$\mathcal{L} = \text{BCE}_w(\sigma(z), y) + \lambda \frac{1}{2|E|} \sum_{(u,v) \in E} p_u p_v, \quad \lambda = 2.$$

Labels come in two flavours. For the small graphs we compute *multi-solution* targets: Gurobi enumerates up to 16 strictly optimal independent sets per graph, and the target y_v is the fraction of those optima that contain node v . Such a target is calibrated, so its own cross-entropy is strictly positive (about 0.19 at $n=50$). For the large benchmarks we use a single near-optimal labelling from the KaMIS solver together with a positive-class weight w ($w=1$ for the small runs, $w=2$ for SATLIB, $w=5$ for the dense ER graphs). Optimisation uses AdamW (weight decay 0.05), a OneCycle learning-rate schedule (peak 3×10^{-4} , 10% warm-up, cosine decay), **bf16** precision, and dropout 0.15, with early stopping on the held-out validation loss.

What the numbers mean.

- **Iteration.** One gradient update on one mini-batch of graphs; “29k iterations” means 29,000 updates. We evaluate on the validation set every 500 iterations and keep the best checkpoint, so runs stop at different counts.
- **Approximation ratio** (pp-AR). The main quality metric. We decode the probabilities greedily (take nodes by probability, skip any neighbour of a chosen node) and report $|S_{\text{ours}}|/|S^*|$; 1.0 is optimal.
- **Feasibility.** The share of decoded solutions that are valid independent sets. Greedy decoding can never pick two adjacent nodes, so it is 1.0 by construction; the difficulty is reaching the optimum, not staying valid.
- **BCE.** Cross-entropy between the predicted probability and the training label. It measures calibration against one labelling, which is not the same as the quality of the final set (see below).
- **Confusion matrix.** Per node, whether predicted MIS membership matches the ground truth, split into true/false positives and negatives.

2 Experiments

We run two studies. The first is a controlled scaling study on Erdős–Rényi (ER) graphs of fixed average degree 7.35 at sizes $n \in \{50, 100, 200, 300\}$, each with a held-out validation split (disjoint seeds) and exact Gurobi labels. The second uses two larger benchmarks from the DIFUSCO suite: SATLIB (SAT formulas turned into conflict graphs, around 1400 nodes) and ER-700-800 (dense ER graphs of 700 to 800 nodes, average degree about 112), both labelled by KaMIS. Every number below is measured on these held-out splits, never on training graphs (for the benchmarks this is the official DIFUSCO test split). Because the same split also drives early stopping and checkpoint selection, it serves as a validation set, so the reported best values are mildly optimistic. The splits hold 5000/5000/1000/300 graphs for $n=50/100/200/300$ and 500 graphs for each benchmark.

2.1 Scaling over problem size

On all four small sizes the model learns cleanly (Fig. 1). The validation BCE drops from its random-init value near 0.73 to between 0.23 and 0.32, while pp-AR climbs from about 0.886 at initialisation to at least 0.98. Larger graphs reach their best checkpoint in fewer iterations because each graph carries more signal, but they also start to overfit sooner since their training pools are smaller (Table 1).




Figure 1: Held-out validation curves for the four small ER sizes. Left: validation BCE (lower is better). Right: approximation ratio pp-AR (higher is better).

Table 1: Main results. The top block is the small-ER scaling study; the bottom block is the two large benchmarks. q is the average MIS fraction $|S^*|/n$. “best BCE” is the lowest validation BCE over the run and “best pp-AR” the highest validation ratio; the two can occur at different checkpoints. Feasibility is 1.00 everywhere thanks to greedy decoding.

Dataset	nodes	train graphs	q	best BCE	best pp-AR	feas.	iter @ best pp-AR
ER $d=7.35$	50	50,000	0.32	0.226	0.982	1.00	29.5k
ER $d=7.35$	100	50,000	0.34	0.268	0.993	1.00	25.5k
ER $d=7.35$	200	12,000	0.36	0.319	0.995	1.00	11.5k
ER $d=7.35$	300	3,000	0.36	0.307	0.997	1.00	15.5k
SATLIB	~1400	9,000	0.33	0.713	0.998	1.00	150k
ER-700-800	700–800	25,000	0.06	n/a	0.903	1.00	325k

2.2 Why a high BCE can still give a high approximation ratio

A striking pattern in Table 1 is that SATLIB reaches the best approximation ratio of the whole study (pp-AR = 0.998) while its BCE stays around 0.7 to 0.8, far above the small ER runs that sit near 0.25. The two metrics measure different things. BCE asks whether each node probability matches one fixed labelling. The benchmark graphs have many different optimal sets, but KaMIS returns only one of them, so a node that belongs to some other optimum is labelled 0 even though the model sensibly gives it a moderate probability. That mismatch, together with the positive-class weight, keeps BCE high. The approximation ratio, by contrast, depends only on the *order* of the probabilities that the greedy decoder reads off: as long as good MIS nodes rank above the nodes that conflict with them, the decoded set is near optimal regardless of absolute calibration. The small ER runs avoid the mismatch because their multi-solution targets are calibrated, so they reach both a low BCE and a high ratio. BCE therefore tracks calibration against one labelling while pp-AR tracks set quality, and only the latter reflects the objective we care about.

2.3 Confusion matrices

Table 2 gives the node-level confusion matrix at each best-pp-AR checkpoint, as a percentage of all nodes. The small ER models are balanced, with roughly equal false positives and false negatives and about 90% node accuracy. The benchmarks behave differently: both over-predict membership (large false-positive share), yet the greedy decoder still recovers near-optimal sets, which is the same decoupling discussed above. The dense ER case is the extreme: only 6% of nodes are in the MIS, the classifier floods the graph with false positives, and node accuracy looks fine

(73%) only because most nodes are correctly rejected.

Table 2: Node-level confusion matrices at the best-pp-AR checkpoint (% of all nodes). TP/FP/FN/TN sum to 100; accuracy is TP + TN.

Dataset	TP	FP	FN	TN	accuracy
ER $n=50$	27.4	4.9	5.0	62.7	90.1
ER $n=100$	28.6	6.1	5.6	59.8	88.4
ER $n=200$	29.9	7.2	5.9	57.0	86.9
ER $n=300$	30.6	6.8	5.8	56.8	87.4
SATLIB	25.6	20.6	7.7	46.1	71.7
ER-700-800	3.4	24.7	2.6	69.3	72.7

2.4 Scaling up, and where it stopped working

The benchmark runs used a 9000-graph training pool. On SATLIB this succeeded and gave the best result of the study, but the run also shows a clear failure mode of prolonged training (Fig. 2). The approximation ratio peaks near 0.998 around iteration 150k and then slowly falls back to 0.95, while the BCE climbs from 0.7 to 1.8. The model keeps fitting the single KaMIS labelling more aggressively, which hurts both calibration and, eventually, the decoded sets. Early stopping on the validation score is therefore essential at this scale.




Figure 2: SATLIB with a 9000-graph pool. The ratio peaks near iteration 150k and then degrades as the BCE rises, a clear overfitting signature.

The same 9000-graph setup did not even start on the dense ER-700-800 distribution: both attempts ran out of GPU memory on a 31.4 GB card inside the GIN message-passing step, before a single evaluation. The cause is the edge count, not the number of graphs. A dense ER graph of 750 nodes has roughly 40 times as many edges per node as a SATLIB conflict graph, so a batch of them produces a far larger message-passing tensor. Whether a run fits depends on how many dense graphs share one GPU batch, not on the pool size: the completed result on this distribution used a larger pool (25,000 graphs) but only 8 graphs per batch instead of 16, and reaches pp-AR = 0.903, well above a degree-heuristic baseline of 0.60 but clearly the hardest case we tried.

3 Key insights

1. **The tiny recurrent network does learn MIS.** Every small-ER size reaches pp-AR between 0.98 and 0.997 on held-out graphs with a single 1.55 M model unrolled for 12 steps, which supports the TRM premise that repeated computation can stand in for scale.
2. **Set quality and node loss are different objectives.** A model can have a poor BCE (SATLIB at 0.7 to 0.8, dense ER far worse) and still decode near-optimal sets, because the greedy decoder only needs the right ranking. Choosing or stopping a run on BCE alone would be misleading; the approximation ratio is the metric to watch.
3. **Density is the real difficulty axis, not size.** The largest sparse graphs (SATLIB, around 1400 nodes) were easier than smaller but much denser ER graphs. Both quality and trainability, including the out-of-memory failure, track edge count and the shrinking MIS fraction q rather than node count.
4. **Feasibility is free, optimality is the work.** Greedy decoding keeps feasibility at 1.00 everywhere, so the feasibility loss mainly shapes the ranking and all remaining effort goes into closing the gap to the optimum.
5. **Long training can hurt.** On the large single-solution benchmark the approximation ratio peaked and then declined with more training, so a held-out stopping criterion matters as much as the model itself.

Reproducibility. All training, data-generation, and plotting code, the dataset specifications, and the configuration files are in the repository at <https://github.com/73gon/trm-mis>.